


OpenACC vs. OpenMP4 P2Z Performance Comparison

Seyong Lee

Oak Ridge National Laboratory

November 23, 2021

P2Z Performance: OpenACC, OpenMP4, and CUDA Performance on a Summit Node





- **ACCV1**: OpenACC Version 1 – Synchronous
- **OMP4V1**: OpenMP4 Version 1 – Synchronous
- **ACCV1-prepin**: OpenACC Version 1 – Synchronous & prepinning host memory
- **OMP4V1-prepin**: OpenMP4 Version 1 – Synchronous & prepinning host memory
- **ACCV3-S1**: OpenACC Version 3 – Asynchronous with 1 stream
- **OMP4V3-S1**: OpenMP4 Version 3 – Asynchronous with 1 stream
- **ACCV3-S1-prepin**: OpenACC Version 3 – Asynchronous with 1 stream & prepinning host memory
- **OMP4V3-S1-prepin**: OpenMP4 Version 3 – Asynchronous with 1 stream & prepinning host memory
- **CUDAV3-S1**: CUDA Version 3 – Asynchronous with 1 stream
- **ACCV3-S10**: OpenACC Version 3 – Asynchronous with 10 streams
- **OMP4V3-S10**: OpenMP4 Version 3 – Asynchronous with 10 streams
- **ACCV3-S10-prepin**: OpenACC Version 3 – Asynchronous with 10 streams & prepinning host memory
- **OMP4V3-S10-prepin**: OpenMP4 Version 3 – Asynchronous with 10 streams & prepinning host memory
- **CUDAV3-S10**: CUDA Version 3 – Asynchronous with 10 streams

V1: Sync version

V3: Async version

P2Z Performance: OpenACC vs. OpenMP4 Performance on an AMD Node

(AMD EPIC 7272 CPU & AMD Instinct MI100 GPU)




- **ACCV1**: OpenACC Version 1 – Synchronous
- **OMP4V1**: OpenMP4 Version 1 – Synchronous
- **ACCV1-prepin**: OpenACC Version 1 – Synchronous & prepinned host memory
- **OMP4V1-prepin**: OpenMP4 Version 1 – Synchronous & prepinned host memory
- **ACCV3-S1**: OpenACC Version 3 – Asynchronous with 1 stream
- **OMP4V3-S1**: OpenMP4 Version 3 – Asynchronous with 1 stream
- **ACCV3-S1-prepin**: OpenACC Version 3 – Asynchronous with 1 stream & prepinned host memory
- **OMP4V3-S1-prepin**: OpenMP4 Version 3 – Asynchronous with 1 stream & prepinned host memory
- **CUDA V3-S1**: CUDA Version 3 – Asynchronous with 1 stream
- **ACCV3-S10**: OpenACC Version 3 – Asynchronous with 10 streams
- **OMP4V3-S10**: OpenMP4 Version 3 – Asynchronous with 10 streams
- **ACCV3-S10-prepin**: OpenACC Version 3 – Asynchronous with 10 streams & prepinned host memory
- **OMP4V3-S10-prepin**: OpenMP4 Version 3 – Asynchronous with 10 streams & prepinned host memory
- **CUDA V3-S10**: CUDA Version 3 – Asynchronous with 10 streams

V1: Sync version



V3: Async version

P2Z Performance: NVIDIA V100 vs. AMD MI100

V100 Peak FP32 Performance: 15.7 TFLOPs



MI100 Peak FP32 Performance: 23.1 TFLOPs




V1: Sync version

V3: Async version

P2Z OpenACC V1 Execution Time Breakdown

OpenACC V1 is a synchronous version



Summit Node (Power 9 CPU + V100 GPU): V100 FP32 Peak Performance = 15.7 TFLOPs

CPU-GPU connection: NVLINK (50 GB/s)



AMD Node (EPIC7272 CPU + MI100 GPU): MI100 FP32 Peak Performance = 23.1 TFLOPs

CPU-GPU connection: PCI-e x16 (31.5 GB/s)

Memory Transfer Performance Comparison

Total H2D Transfer Size: 3.8 GB

Total D2H Transfer Size: 0.525 GB



Normalized against the peak communication bandwidth

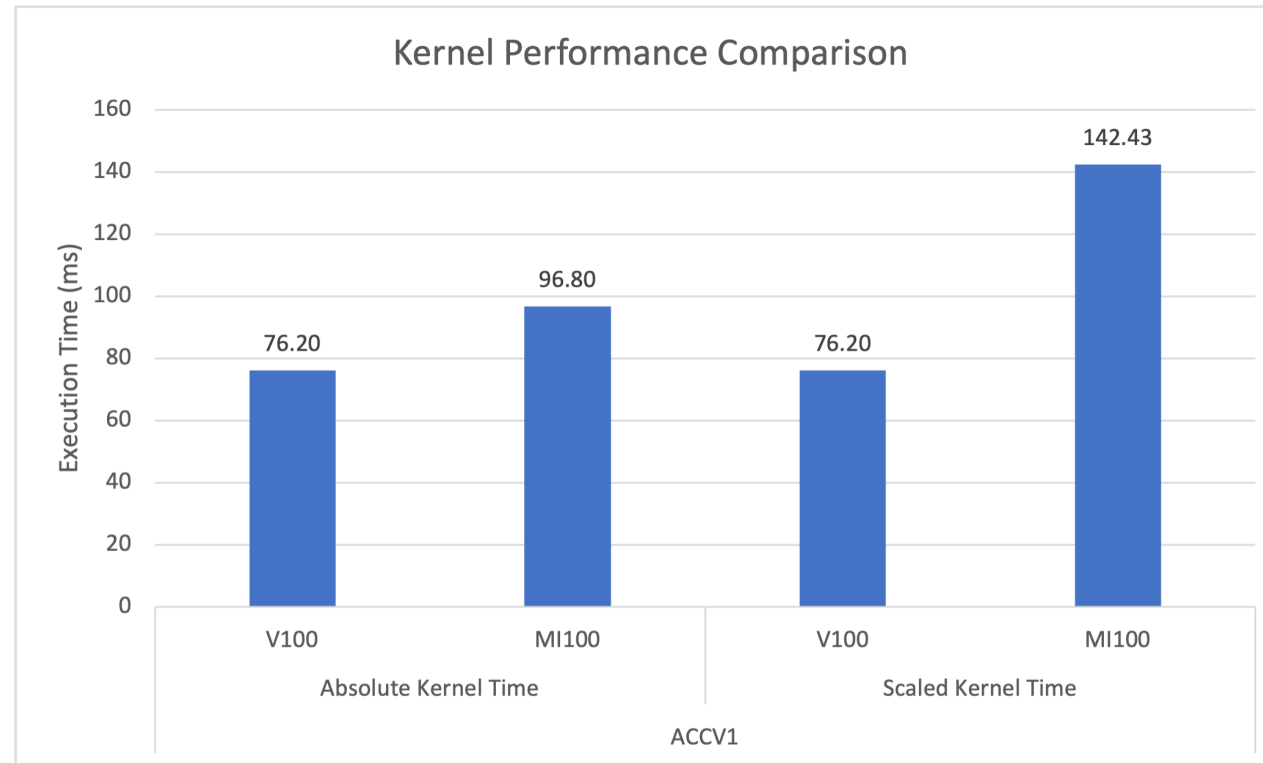
Summit Node (Power 9 CPU + V100 GPU): V100 FP32 Peak Performance = 15.7 TFLOPs

CPU-GPU connection: NVLINK (50 GB/s)

AMD Node (EPIC7272 CPU + MI100 GPU): MI100 FP32 Peak Performance = 23.1 TFLOPs

CPU-GPU connection: PCI-e x16 (31.5 GB/s)

Kernel Performance Comparison



Scaled kernel time of MI100 is the estimation when assuming MI100 has the same FP32 peak performance as V100 (15.7 TFLOPs).

Summit Node (Power 9 CPU + V100 GPU): V100 FP32 Peak Performance = 15.7 TFLOPs

CPU-GPU connection: NVLINK (50 GB/s)

AMD Node (EPIC7272 CPU + MI100 GPU): MI100 FP32 Peak Performance = 23.1 TFLOPs

CPU-GPU connection: PCI-e x16 (31.5 GB/s)