P2Z CUDA Version Performance on Summit

Seyong Lee

Oak Ridge National Laboratory

October 18, 2022

P2Z CUDA Version Performance on a Summit Node (NITER = 10)

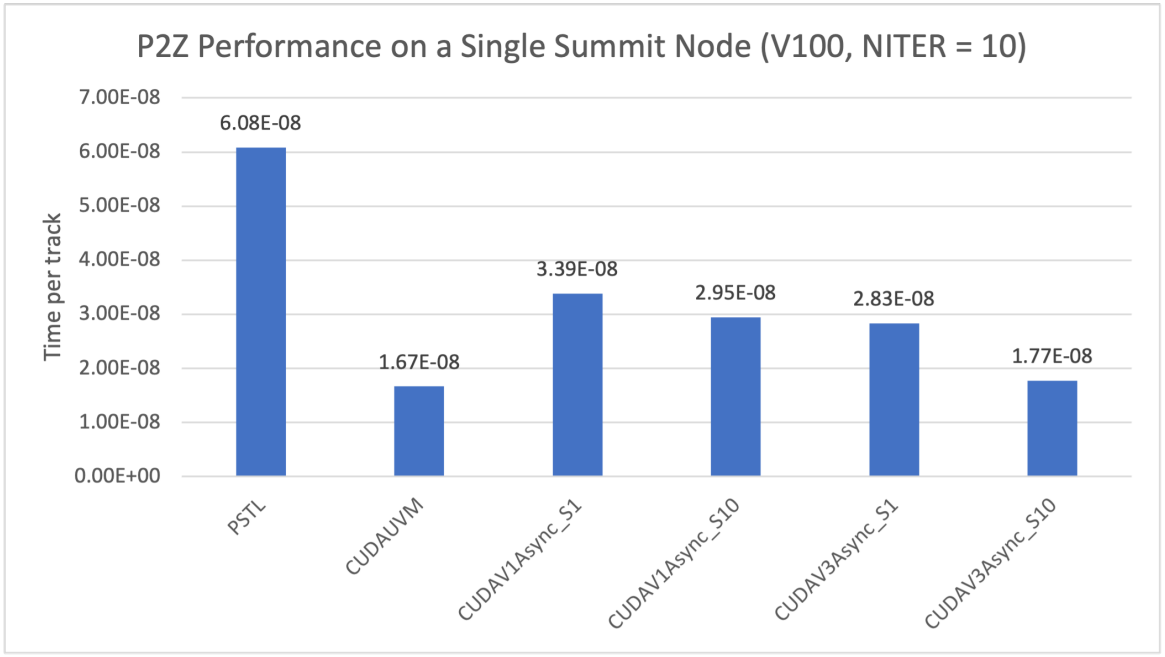

Measured time includes both memory transfers and computations

- **PSTL**: C++ parallel STL
- **CUDAUVM**: CUDA unified memory version (async, number of streams = 1)
- **CUDAV1Async_S1**: CUDA V1 async unified memory version: number of streams = 1
- **CUDAV1Async_S10**: CUDA V1 async unified memory version: number of streams = 10
- **CUDAV3Async_S1**: CUDA V3 async version: number of streams = 1
- **CUDAV3Async_S10**: CUDA V3 async version: number of streams = 10

Compilers: NVCC (V11.0)

P2Z CUDA Version Performance on a Summit Node (NITER = 1)

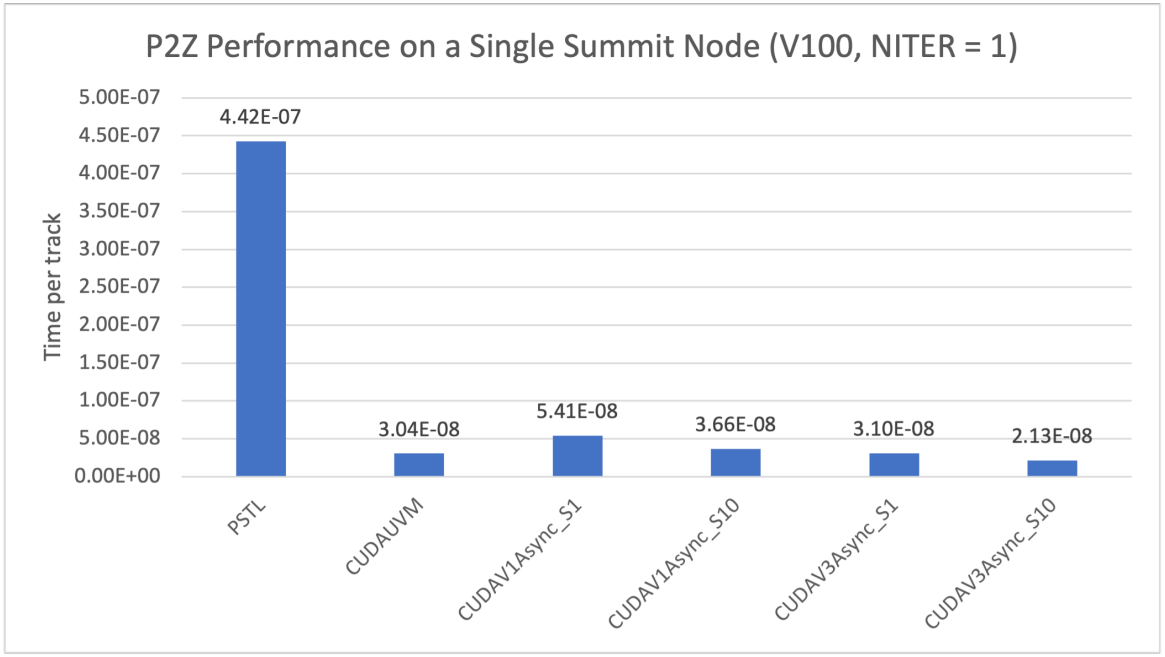

Measured time includes both memory transfers and computations

- **PSTL**: C++ parallel STL
- **CUDAUVM**: CUDA unified memory version (async, number of streams = 1)
- **CUDAV1Async_S1**: CUDA V1 async unified memory version: number of streams = 1
- **CUDAV1Async_S10**: CUDA V1 async unified memory version: number of streams = 10
- **CUDAV3Async_S1**: CUDA V3 async version: number of streams = 1
- **CUDAV3Async_S10**: CUDA V3 async version: number of streams = 10

Compilers: NVCC (V11.0)